

Codernize AI – Design Document

Andrea Paparella

May 19, 2025

Codernize AI is a set of Python scripts representing three different tools designed to generate documentation, diagrams, and modernization reports for a given repository. The tool is currently built for Java projects but can be extended to support other project types.

1 DocGen: Documentation Generation Tool

1.1 Introduction

DocGen automates the creation of documentation by scanning a repository for source and configuration files, then generating a single document describing the project.

1.2 Pipeline Evolution

Multiple iterations were explored before reaching the current version of the tool that delivers reliable and readable documentation.

v0 - One File per Doc

- Walk repository; generate one documentation file per source file.
- *Outcome*: Huge scatter of micro-docs; navigation was impractical.

v1 - Single Growing Document

- Append each new file's doc to a single Markdown file as the scan proceeds.
- *Issue*: Recency bias, later files dominated the context window, so earlier sections were forgotten or overwritten.

v2 - Current Architecture

1. Scan repository for relevant extensions.
2. Categorize each file.
3. Generate a concise snippet per file.
4. Aggregate all snippets in one shot into a unified draft.
5. Polish the draft with a second model pass focused on layout, headings, and flow.

High-Level Architecture Design

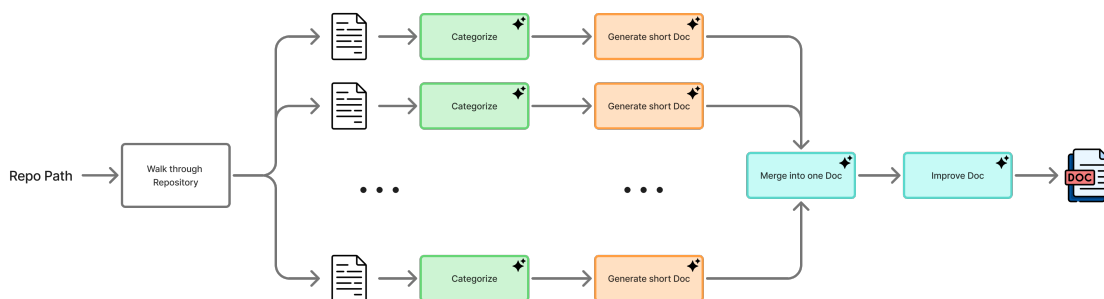


Figure 1: High-level architecture design of DocGen-v2.

Design Choices: Output Format

Markdown was chosen over AsciiDoc because, despite AsciiDoc’s richer syntax, the LLM often generated inconsistent or error-prone output with it. In contrast, Markdown consistently produced well-formatted results.

1.3 Model Selection Strategy

Task	Model	Why
File categorisation	<code>gpt-4.1-nano</code>	Fast, low-latency; prompt is highly constrained (single-word category).
Per-file short docs	<code>gpt-4.1-nano</code>	Output is short and templated; speed outweighs marginal quality gains of a larger model.
Combined summary	<code>gpt-4o-mini</code>	Needs larger context window and stronger reasoning to fuse many snippets coherently.
Final cleanup / re-org	<code>gpt-4o-mini</code>	Same reasons as above; excels at editing and structural refinements.

1.4 Running the Tool

Set-up

Before running the script, export your OpenAI API key:

```
export OPENAI_API_KEY = <your_openai_api_key>
```

Usage

```
python3 gen_docs.py /path/to/project -o docs -doc docs.md -d
```

- `/path/to/project` – Root directory of the target project.
- `-o docs` – Output directory for generated documentation.
- `-doc docs.md` – Name of the final documentation file.
- `-d` – Enables debug mode (detailed logging and per-file documentation snippets).

2 ModGen: Modernization Report Generation Tool

2.1 Introduction

ModGen generates a modernization report by scanning source and configuration files, identifying modernization opportunities, and suggesting boilerplate code for new architectures (e.g., Spring Boot). To improve understanding of the target repository, the tool also accepts project documentation as input. You can use documentation generated by DocGen for this purpose.

2.2 Pipeline Evolution

Fewer iterations were needed thanks to lessons learned from DocGen. Instead of providing the entire repository to the LLM as input, each file is processed separately, and the outputs are then merged, similar to the approach used in DocGen.

v0 – Current Architecture

1. Scan repository for relevant file extensions.
2. Generate a modernization analysis per file, suggesting Spring Boot migration opportunities.
3. Aggregate the analyses into a unified report, using the repository documentation as additional input to give more comprehensive context to the LLM.

High-Level Architecture Design

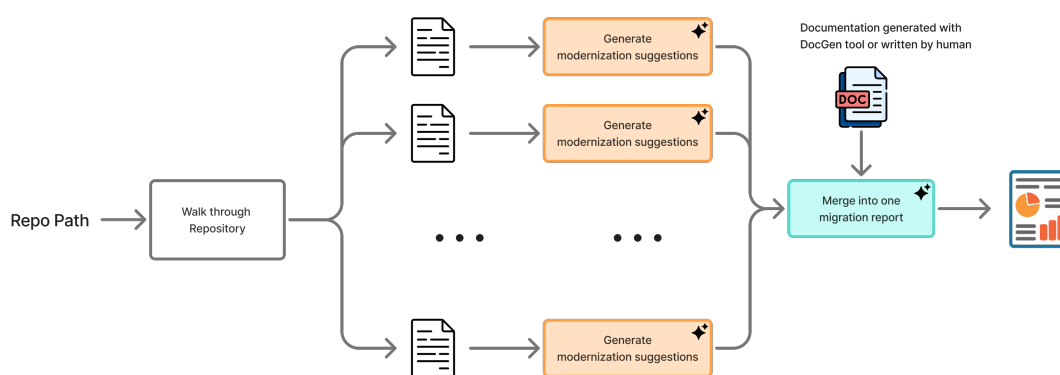


Figure 2: High-level architecture design of ModGen-v0.

2.3 Model Selection Strategy

Task	Model	Justification
Per-file modernization analysis	<code>gpt-4.1-mini</code>	High reasoning capability and large context window to produce detailed analyses and high-quality code.
Combined report	<code>gpt-4.1-mini</code>	Same reasoning as above; suited for integrating comprehensive insights.

2.4 Running the Tool

Setup

Before running, export your OpenAI API key:

```
export OPENAI_API_KEY=<your_openai_api_key>
```

Usage

```
python3 mod-gen.py /path/to/docs.md /path/to/project -o docs -report  
modernization-report.md -d
```

- `/path/to/docs.md` – Documentation path to give the LLM better context.
- `/path/to/project` – Root directory of the target project.
- `-o docs` – Output directory for the generated documentation.
- `-report modernization-report.md` – Name of the generated modernization report.
- `-d` – Enables debug mode (detailed logging).

3 DiagGen: Diagram Generation Tool

3.1 Introduction

DiagGen generates two diagrams of a system: a complete version and a simplified version, by scanning source and configuration files.

3.2 Pipeline Evolution

Fewer iterations were needed due to learnings from DocGen. Instead of providing the entire repository to the LLM as input, each file is processed separately, and the outputs are then merged, similar to the approach used in DocGen.

3.2.1 v0 - Single Diagram as Output

- Walk the repository; generate one diagram per file.
- Aggregate all diagrams in one shot into a unified diagram.
- *Outcome*: The resulting diagram was overly complex.

3.2.2 v1 - Two Diagrams as Output

To avoid having a single complex diagram that might be hard to read and understand, the pipeline also generates a simplified version from the more detailed one.

1. Scan the repository for relevant file extensions.
2. Generate a diagram per file.
3. Aggregate all diagrams into a unified diagram; store it.
4. Simplify the diagram, retaining only the most important information.

High-Level Architecture Design

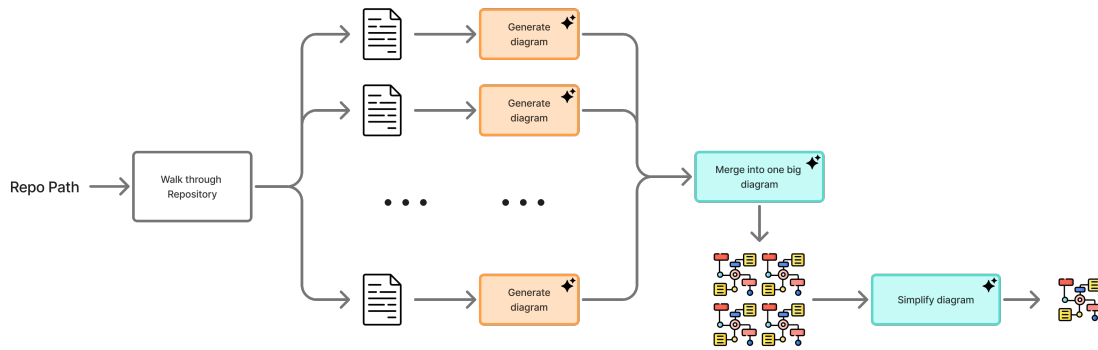


Figure 3: High-level architecture design of DiagGen-v1.

3.3 Model Selection Strategy

Task	Model	Justification
Per-file diagram	<code>gpt-4.1-nano</code>	Good reasoning capability and fast execution.
Combined diagram	<code>gpt-4.1-mini</code>	High reasoning ability; suitable for integrating comprehensive insights.
Simplified diagram	<code>gpt-4.1-mini</code>	Same reasons as above.

3.4 Running the Tool

Setup

Before running, export your OpenAI API key:

```
export OPENAI_API_KEY=<your_openai_api_key>
```

Usage

```
python3 diag-gen.py /path/to/project -o docs -diagram system-diagram.mermaid -d
```

- `/path/to/project` – Root directory of the target project.
- `-o out_docs` – Output directory for generated documentation.
- `-diagram system-diagram.mermaid` – Name of the final documentation file.
- `-d` – Enables debug mode (detailed logging).

***Note:** Diagram generation may produce invalid output that fails to render in Mermaid. You may need to run it multiple times or manually fix the result.*

4 Future Improvements

- **Extension handling:** The tools currently rely on a hardcoded list of file extensions to identify relevant files (e.g., `.java`, `.html`). This approach is not scalable. One possible improvement is to scan all file extensions in the repository and use an LLM to determine which ones should be included in documentation generation.
- **Handling large files and token limits:** If a file is too large, or if the combined content of multiple files exceeds the model's context window, the API may fail and skip the file. To address this, we could explore tools like LangChain to chunk large files or summarize groups of files before generating documentation. However, this may lead to information loss. Alternatively, switching to models with larger context windows can help, although they still have upper limits.
- **Parallelization of per-file generation:** Currently, the tool generates documentation or analyses for each file sequentially. This could be optimized by running multiple model calls in parallel, reducing runtime. The aggregation step could then proceed once all per-file snippets are available.
- **Diagram generation reliability:** In some cases, the generated diagrams contain syntax errors or malformed content that prevent Mermaid renderers from displaying them correctly. A possible future improvement would be to fine-tune a model specifically for diagram generation tasks.